

# THREADS

---

## REFERENCES:

1. “OPERATING SYSTEM CONCEPTS” 9<sup>TH</sup> EDITION BY ABRAHAM SILBERSCHATZ, PETER BAER GALVIN AND GREG GAGNE
2. “OPERATING SYSTEMS: INTERNALS AND DESIGN PRINCIPLES”, 7TH EDITION BY WILLIAM STALLINGS
3. “INTRODUCTION TO OPERATING SYSTEMS ”, BY PROF. CHESTER REBEIRO  
([HTTPS://NPTEL.AC.IN/COURSES/106/106/106106144/](https://npTEL.AC.IN/COURSES/106/106/106106144/))

## CONSIDER THE FOLLOWING SCENARIO:

- There are 4 CPUs in the system
- The given program is adding numbers upto 10 million using an addall() and executing the process in one CPU.

**Problem:** Other processors are not utilized and single process takes long time to complete execution

```
#include <stdio.h>

unsigned long addall(){
    int i=0;
    unsigned long sum=0;

    while (i< 100000000){
        sum += i;
        i++;
    }
    return sum;
}

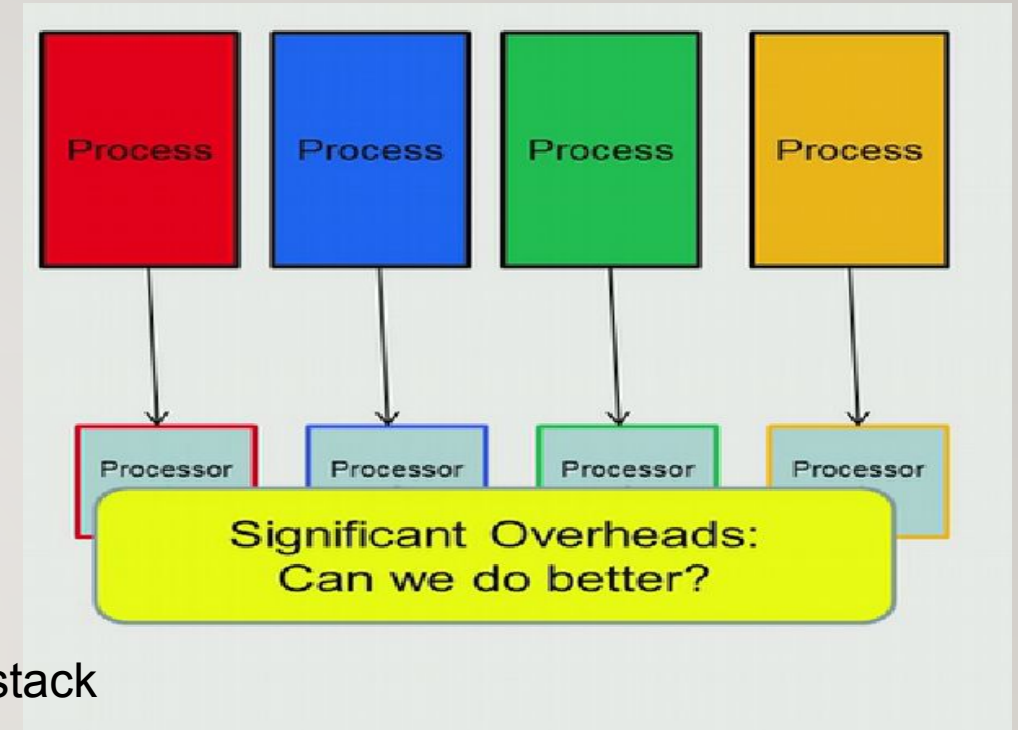
int main()
{
    unsigned long sum;
    srand(time(NULL));
    sum = addall();
    printf("%lu\n", sum);
}
```

## What can be the better method to perform this summation?

- Create 4 processes such that each process adds 2.5 million no.s
- We need 4 fork() calls to create 4 processes
- Each process can execute in one processor
- Reduces the computation time,

### But

- Each process has its own set of instructions, data, heap and the stack
- A large portion of these 4 processes are similar.
- A lot of duplication of instructions and data
- Process Management and IPC required

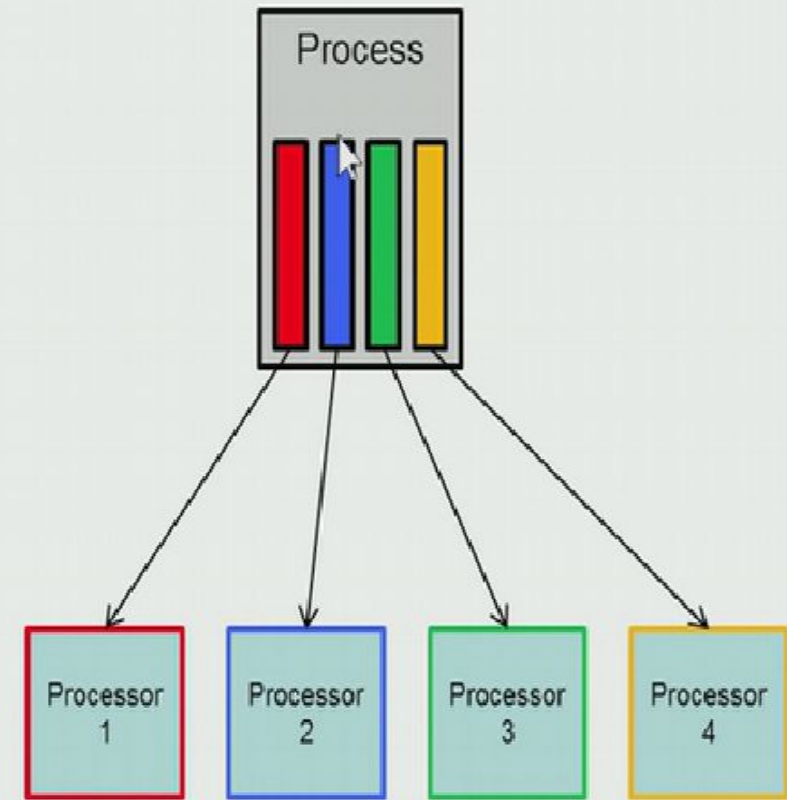


Ref. 3



## Any better method to reduce this overhead and achieve parallelization?

- Create 4 threads under 1 process, using Pthread.
- Each thread executes in separate processor
- Each thread shares common instructions, parameters and heap, etc.
- However, each thread has separate stack.
- Each thread will add 2.5 million no.s
- Threads are lighter than processes
- Very few or no system calls needed to create threads



Ref. 3

# Threads


- Threads are separate streams of execution within a single process.
- They are not isolated from each other.
- The state of a thread is stored in Thread Control Block which contains registers and stack
- It provides mechanisms to perform multiple tasks concurrently.
- Each thread has got associated:

Thread ID

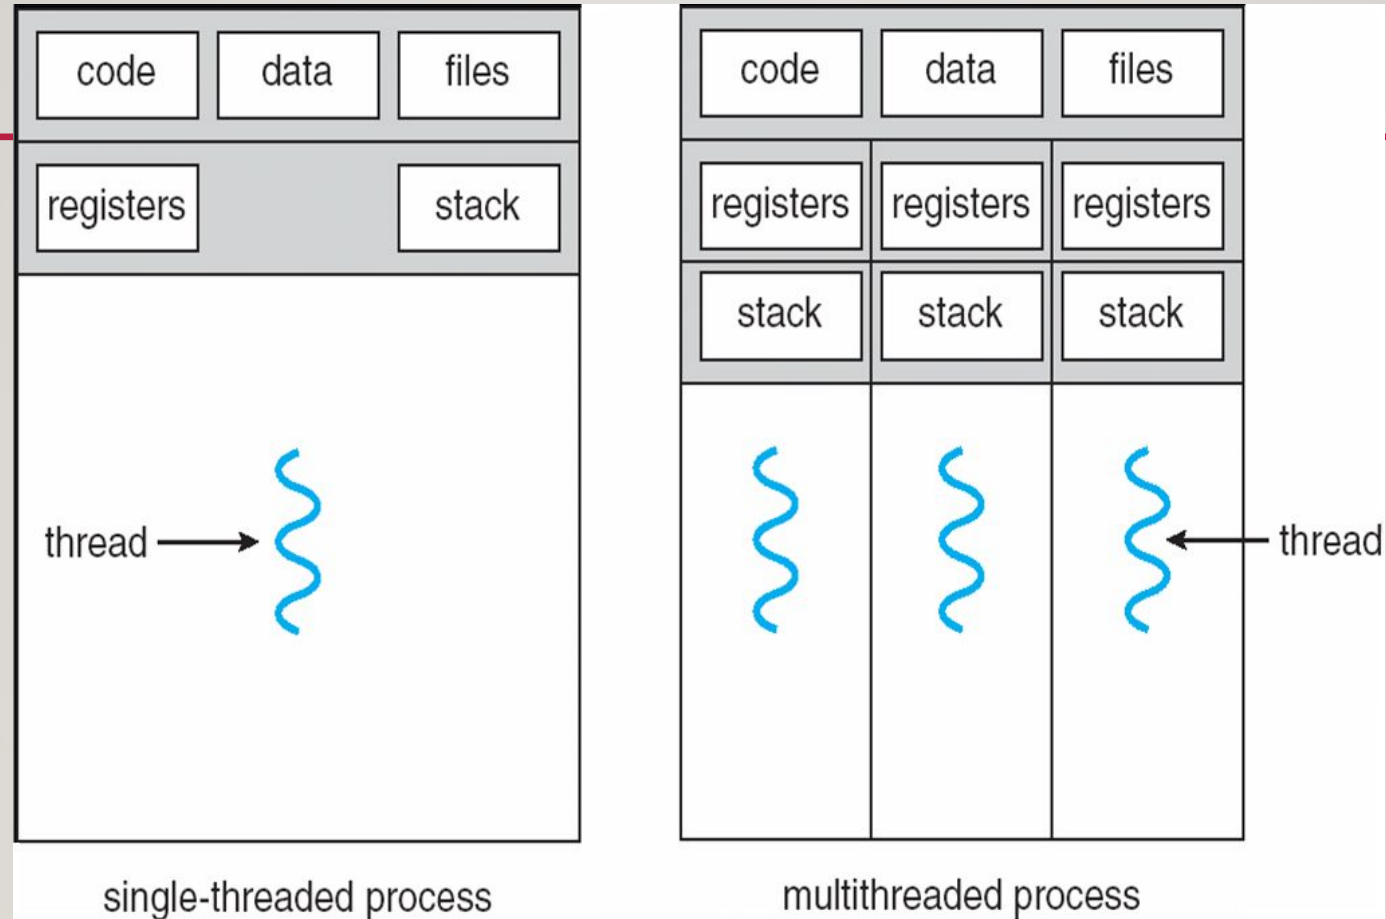
Program counter

Register set

Stack



# SINGLE AND MULTITHREADED PROCESSES



**heavyweight** process

**lightweight** process

Ref. 1



# Threads vs Processes

- A thread has no data segment or heap
- A thread cannot live on its own. It needs to be attached to a process
- There can be more than one thread in a process. Each thread has its own stack
- If a thread dies, its stack is reclaimed
- A process has code, heap, stack, other segments
- A process has at-least one thread.
- Threads within a process share the same code, files.
- If a process dies, all threads die.

Based on Junfeng Yang's lecture slides

<http://www.cs.columbia.edu/~junfeng/13fa-w4118/lectures/I08-thread.pdf>

# MERITS OF USING THREADS

---

- Threads can be created and destroyed quickly as compared to processes
- Applications can use threads to execute some functions in the background
- Threads can share the same address space.
- It takes less time to switch between threads due to smaller state record




# THREADS SCHEDULING

---

- Threads are scheduled to execute in CPU independently
- State of each executing thread is maintained separately.
- If a process is suspended, then all its threads are suspended
- If a process is terminated, then all its threads are terminated
- A thread also has states like ready, running, waiting or blocked.

# TYPES OF THREADS

---



NOTE: We are talking about threads for *user* processes. Both ULT & KLT execute in user mode. An OS may also have threads but that is not what we are discussing here.

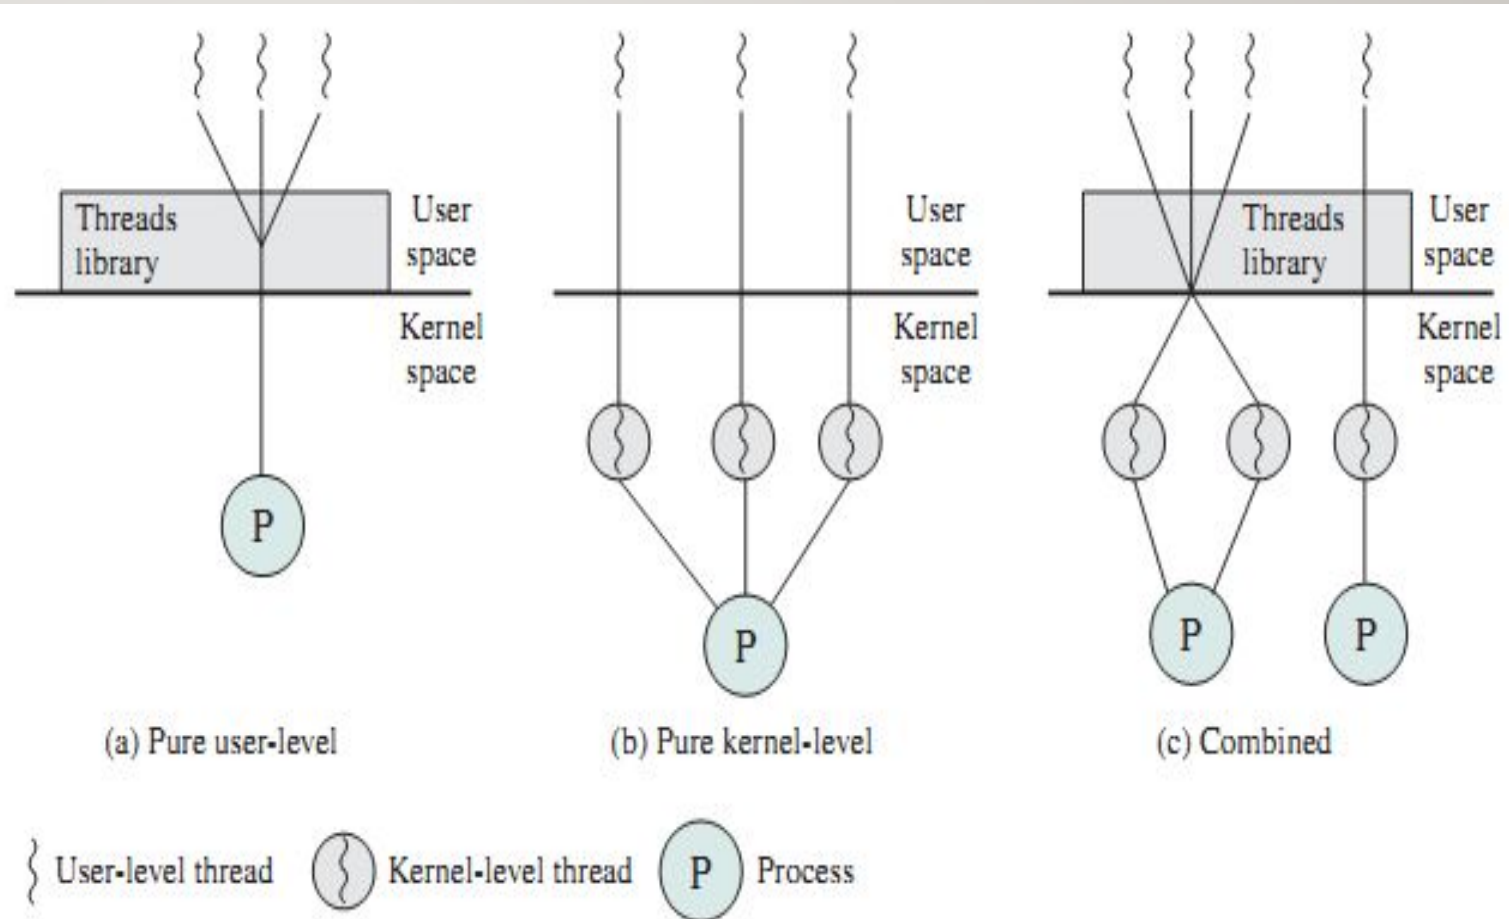
# THREADS MANAGEMENT

## User-Level Threads (ULTs)

- Managed by applications and user level thread library
- Kernel is not aware of these threads

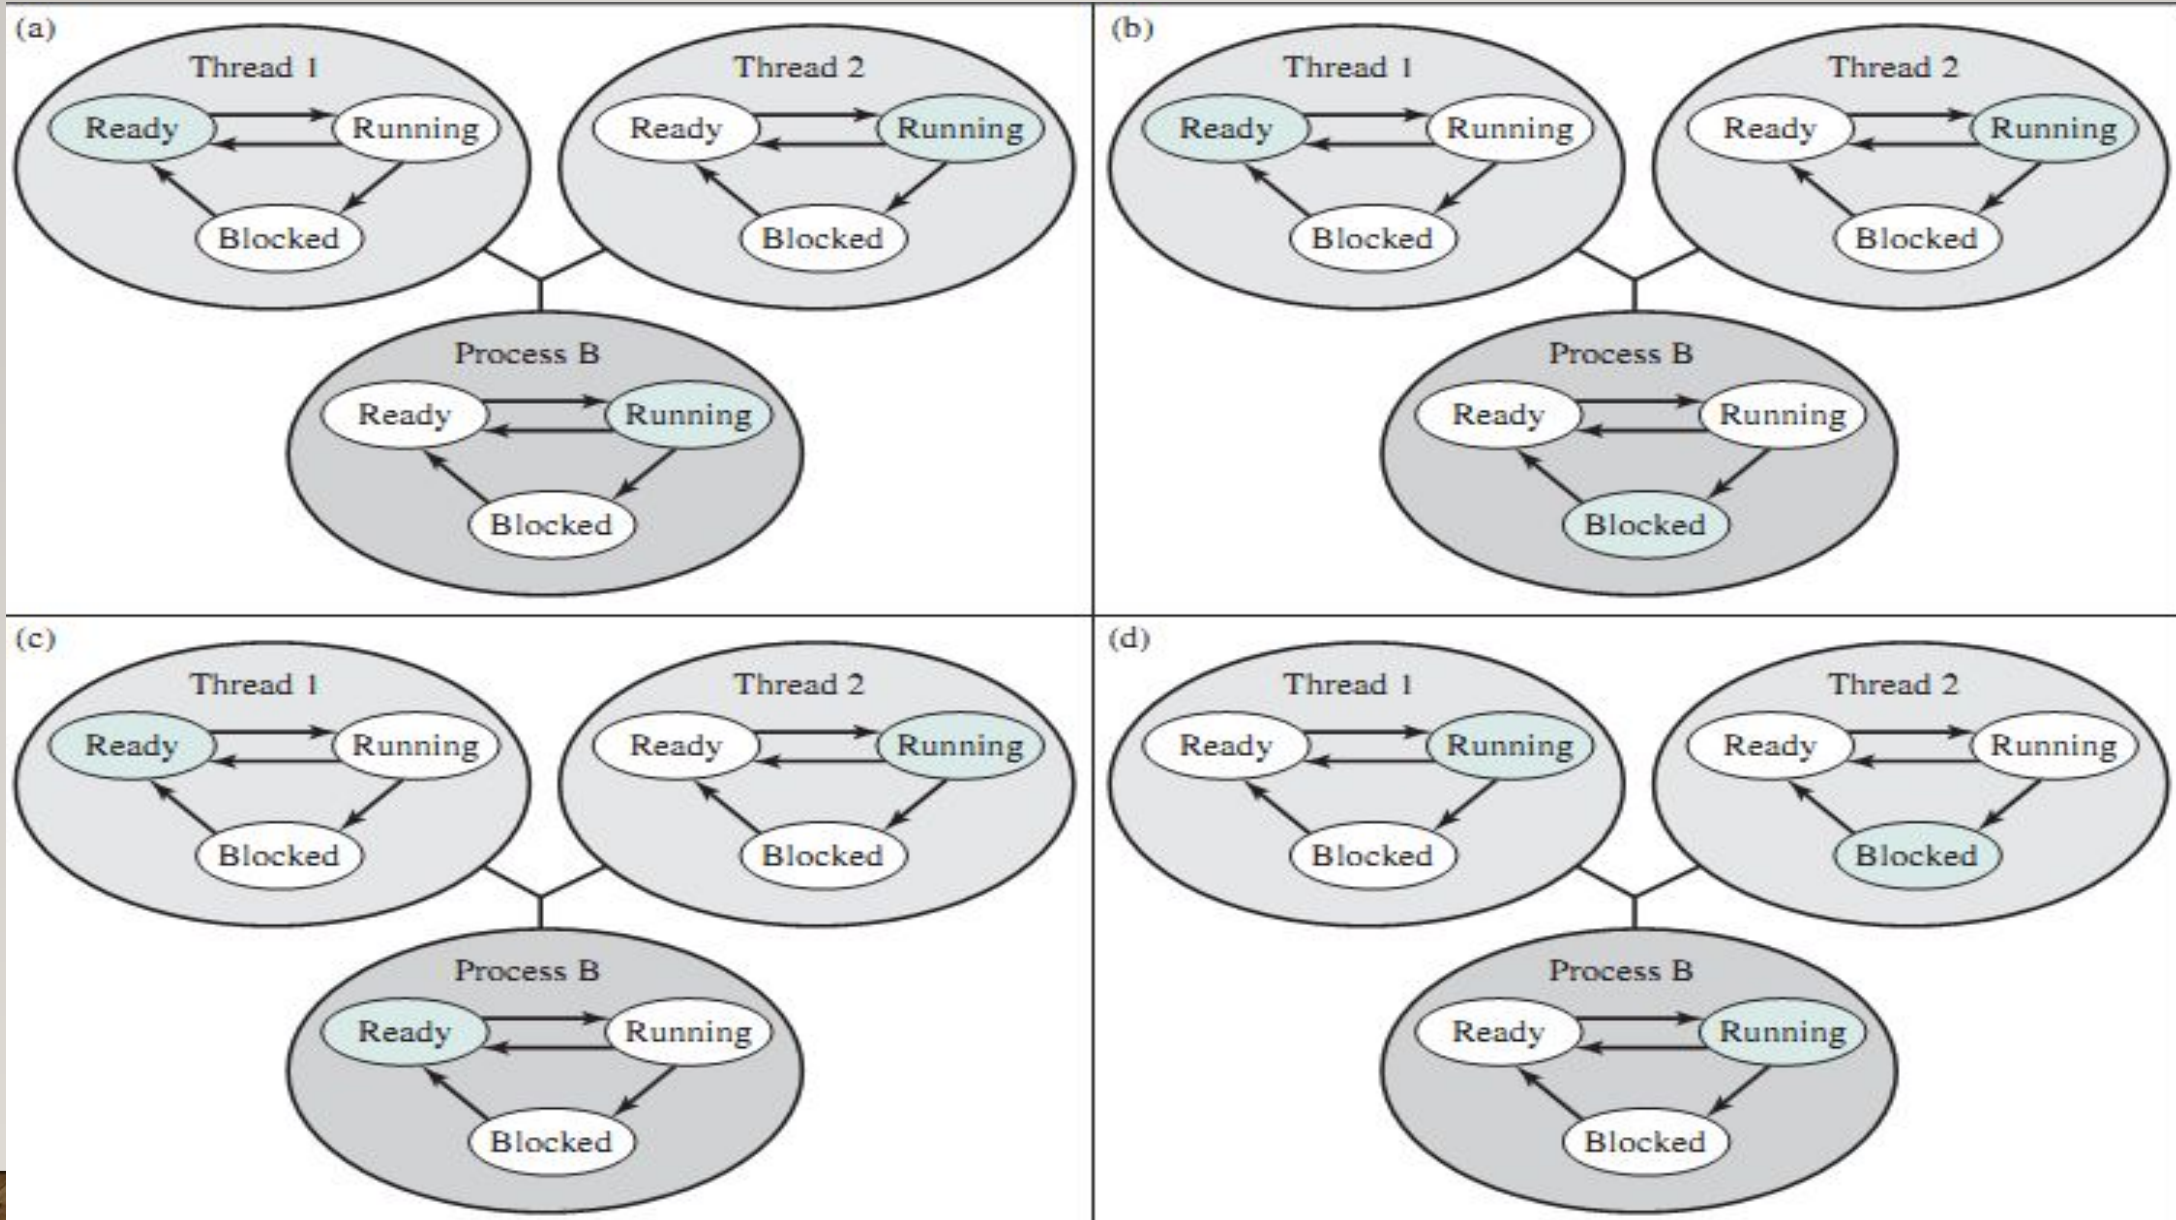
## Kernel-Level Threads (KLTs)

- These are created and managed by kernels
- Also called as light weight process





# Relationships Between ULT States And Process States




# MERITS AND DEMERITS OF ULT

- + It can be implemented in OS that does not support threading
- + Fast creation and switching
- + Does not need System call
- Process with many threads also competes with a single threaded process
- Scheduling decisions cannot be made to favour processes with the larger number of threads
- If one thread makes system call then all others get blocked

Solution: JACKETING


converts a blocking system call into a non-blocking system call



Ref.2

# MERITS AND DEMERITS OF KLT

- + Thread table is stored in kernel space. Hence, kernel knows about no. of threads a process has
  - + OS can provide more time quantum to a process with large no. of threads
  - + Better to use for application that frequently blocks
  - + One thread making system call does not block others
- 
- Slow
  - Larger overhead due to kernel level management
  - The transfer of control from one thread to another within the same process requires a mode switch to the kernel



Ref.2



# MULTITHREADING MODELS


---

- One-to-One
- Many-to-One
- Many-to-Many

# MANY-TO-ONE

---

Many user-level threads mapped to single kernel thread




# ONE-TO-ONE

---

Each user-level thread maps to kernel thread


- Examples
  - Windows NT/XP/2000
  - Linux





# MANY-TO-MANY MODEL

- Many user level threads are mapped to many kernel threads
- It allows the operating system to create a sufficient number of kernel threads
- Example
  - Windows NT/2000



# THREAD LIBRARIES

---

- It provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS
- Three main thread libraries in use today:
  - POSIX Pthreads
  - Win32
  - Java

# POSIX

- It can be used over Linux systems
  - Pthreads API must be compiled with ***-pthread*** or ***-lpthread***
- 

```
#include <pthread.h>
```

```
pthread_t pthread_self()
```

**returns** : ID of current (this) thread



# pthread library

- Create a thread in a process

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine) (void *),  
                  void *arg);
```

Thread identifier (TID) much like

Pointer to a function,  
which starts execution in a  
different thread

Arguments to the function

- Destroying a thread

```
void pthread_exit(void *retval);
```

- Join : Wait for a specific thread to complete

```
int pthread_join(pthread_t thread, void **retval);
```

TID of the thread to wait for

Exit status of the thread

# Example

```
#include <pthread.h>
#include <stdio.h>

unsigned long sum[4];

void *thread_fn(void *arg){
    long id = (long) arg;
    int start = id * 2500000;
    int i=0;

    while(i < 2500000){
        sum[id] += (i + start);
        i++;
    }
    return NULL;
}

int main(){
    pthread_t t1, t2, t3, t4;

    pthread_create(&t1, NULL, thread_fn, (void *)0);
    pthread_create(&t2, NULL, thread_fn, (void *)1);
    pthread_create(&t3, NULL, thread_fn, (void *)2);
    pthread_create(&t4, NULL, thread_fn, (void *)3);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    pthread_join(t3, NULL);
    pthread_join(t4, NULL);
    printf("%lu\n", sum[0] + sum[1] + sum[2] + sum[3]);
    return 0;
}
```

Note. You need to link the pthread library



```
$ gcc threads.c -lpthread
$ ./a.out
```

# TERMINATING THREAD

---

```
#include <pthread.h>
```

```
void pthread_exit (return_value)
```

Threads terminate in one of the following conditions:

- Completes function execution and return value
- pthread\_cancel() request received by thread
- Thread initiates termination
- The process of the threads terminates



# THREAD CANCELLATION

---

- *pthread\_cancel()* : Terminates a thread before it has completed its execution

**Whether thread cancel or not depends in its state and type**

## **States**

- PTHREAD\_CANCEL\_DISABLE: Thread can not be cancelled.
- PTHREAD\_CANCEL\_ENABLE: This is default state. Thread can be cancelled



# THREAD CANCELLATION

---

- Two types of thread cancellation
  - **Asynchronous cancellation:** terminates the target thread immediately
    - PTHREAD\_CANCEL\_ASYNCHRONOUS
  - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
    - PTHREAD\_CANCEL\_DEFERRED: Cancel when thread reaches 'cancellation point'

# THREADING ISSUES

---

- Use of **fork()** and **exec()** system calls
- Signal handling
- Thread pools
- Thread safety
- Thread-specific data

# USE OF *FORK()* ,*EXEC()*, *EXIT()*

- 
- **Does *fork()* duplicate only the calling thread or all threads?**
  - Few unix OS keep two version of fork to have both the options.
  - ***Exec()*:** the program specified in the parameter to *exec()* will replace the entire process—including all threads
  - **Recommendation:** In a process of multiple threads use *fork()* only after *exec()*

# SIGNAL HANDLING

- Signals are used to notify about events to a process.
  - A signal handler is used to process signals in the following way:
    1. Signal is generated by particular event
    2. Signal is delivered to a process
    3. Signal is handled
  - Signal delivery options:
    - To the intended thread
    - To every thread in the intended process
    - To certain threads in the process
    - Assign a specific thread to receive all signals for the process
-



# THREAD POOLS

---

- Create and maintain a number of threads in a pool
- Assign work to the threads as per the need
- Faster method to handle a request using an existing thread instead of creating a new one
- It bounds the number of threads in the application(s) to the size of the pool

# THREAD SAFETY

A function is called *thread-safe* when it can be called by multiple threads at the same time without creating any disruptions.

Example of a function i.e. not safe:

---

```
static int glob = 0;
static void Incr (int loops)
{ int loc, j;
  for (j = 0; j<loops; j++ {
    loc = glob;
    loc++;
    glob = loc;}
}
```

Employs global or static values that are shared by all threads

# HOW TO ENSURE THREAD SAFETY?

---

- Serialize the function: Keep the critical section of the code locked so that only one thread access it at a time other threads out
- Use only thread safe system functions
- Avoid use of global and static variables

# THREAD SPECIFIC DATA

---

- Makes existing functions thread-safe .
  - May be slightly less efficient than being reentrant
- Allows each thread to have its own copy of data
  - Provides per-thread storage for a function
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)



# THREADS: Pros and Cons

---

- Advantages of multithreading
  - Easy to share resources and faster to create
- Disadvantages of multithreading
  - Compete for acquiring memory
  - Ensure threads-safety
  - Error in one can disrupt the execution of other threads due to sharing of resources
- Considerations for future design
  - Handling signals is tricky
  - All threads must run the same program

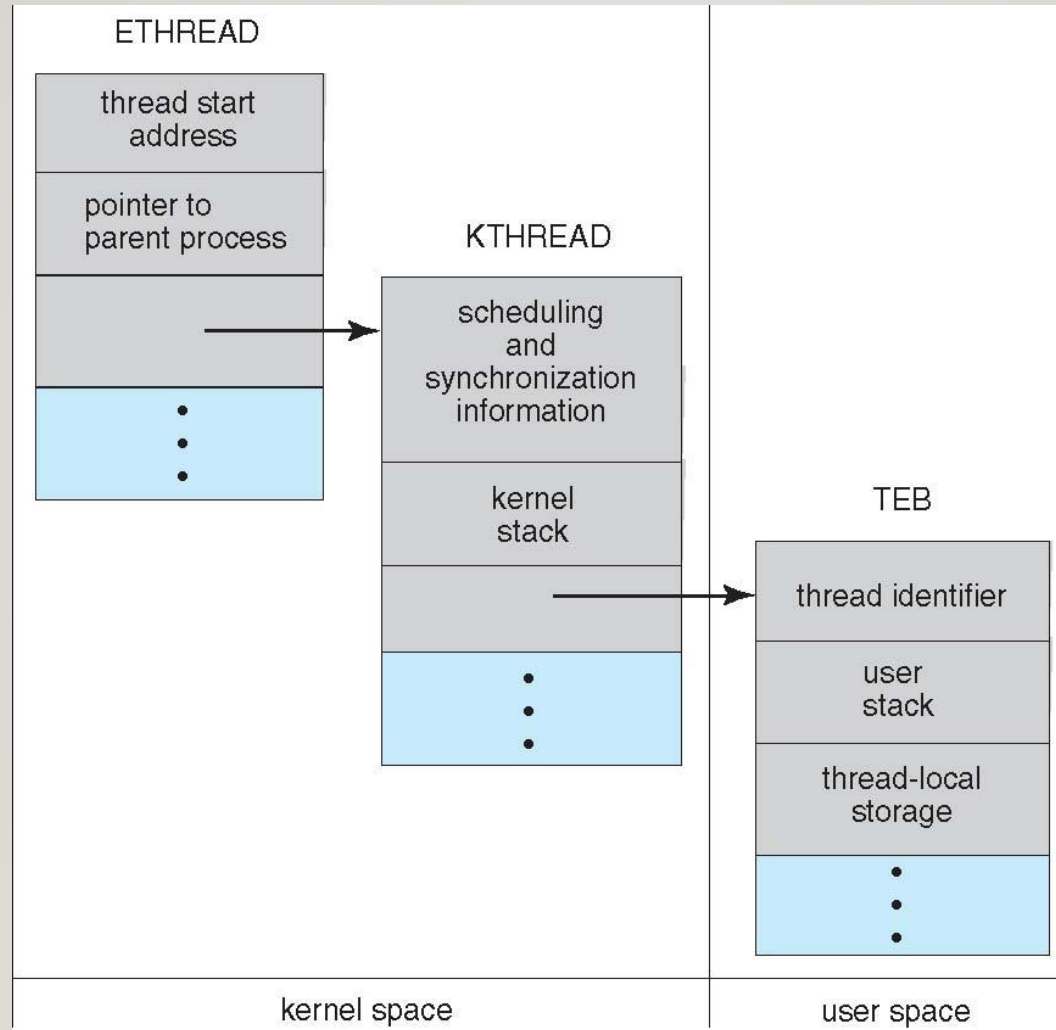
# OPERATING SYSTEM EXAMPLES

- Windows XP Threads
- Linux Thread

# WINDOWS XP THREADS

- It implements one-to-one mapping of threads with kernel-level
- 
- Each thread contains
    - A unique thread id
    - Set of Registers
    - Separate user and kernel stacks
    - Private data storage area
  - These are called context of the threads
  - The primary data structures of a thread include:
    - ETHREAD (executive thread block)
    - KTHREAD (kernel thread block)
    - TEB (thread environment block)

# WINDOWS XP THREADS





# LINUX THREADS

- Threads are referred as Tasks in Linux
- 
- Tasks are created using **clone()** system call
  - **clone()** allows a child task to share the address space of the parent task (process)

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.