

## Logical Clocks

6. (14 points) Three computers at CMU, A, B, and C communicate using a protocol that implements the idea of lamport clocks (they include their clock time stamp in messages). For reference, if you need a reminder, recall that the three rules of Lamport's algorithm are:

1. At process  $i$ , increment  $L_i$  before each event
2. To send message  $m$  at process  $i$ , apply rule 1 and then include the current local time in the message, i.e.,  $\text{send}(m, L_i)$ .
3. To receive a message  $(m, t)$  at process  $j$ , set  $L_j = \max(L_j, t)$  and then apply rule 1 before time-stamping the receive event.

At the beginning of time, all three computers begin with their logical clock set to zero (0). Later, the following sequence of events occurs:

- A sends message M1 to B: "hi".
- After receiving M1, B sends message M2 to C: "A told me hi"
- After receiving M2, C sends message M3 to A: "B is boring"

- (a) Indicate the time included with the messages as they are sent at each step.

$\text{Send}(M1, \_)$   
 $\text{Send}(M2, \_)$   
 $\text{Send}(M3, \_)$

**Solution:**

$\text{Send}(M1, 1)$   
 $\text{Send}(M2, 3)$   
 $\text{Send}(M3, 5)$

- (b) Maintaining all clock states from the previous question, three ADDITIONAL messages are sent:

- After receiving M3, A sends message M4 to B: "C is kind of random!"
- After receiving M4, B sends message M5 to A: "C is boring"
- A receives message M5

After all of these messages have been sent and received, what time does each computer think it is?

A	
B	
C	

**Solution:** A=10, B=9, C=5.

(Remember, receiving and sending are different events!)

- (c) Is this a relatively or totally ordered system?

**Solution:** This is a relatively ordered system.

Reset all clocks to zero. This time:

- A sends message M1 to B.
- The user of machine A is talking on the phone with the user of machine B. She tells him "I just sent you a message online. Please send a message to C right now."
- B sends message M2 to C.

(d) Once all of the messages have been sent and received, what times might C's clock be set to under which circumstance?

**Solution:** If B received the message before sending the message to C, then C's clock would read 4.

Otherwise, C's clock would read 2.

(e) Explain briefly why the scenarios above could happen even though the events must have happened sequentially in real-time.

**Solution:** Because the communication over the cell phones was not timestamped using the lamport clock protocol. Therefore, the protocol did not know that the message from A to B had to have happened before the message from B to C.

(f) Which of your two answers for C's clock is more likely to arise? Briefly (1 sentence) justify why.

**Solution:** It's more likely to be 5. Most of the time, the communication latency between two computers at CMU is likely to be in the few milliseconds range, which is much faster than the two humans can communicate commands in words. Therefore, message M1 is likely to reach machine B well before the user instructs it to send a message to C.

# Homework 1 (Time, Synchronization and Global State) - 100 Points

CS425 Distributed Systems, Fall 2009, Instructor: Klara Nahrstedt

Out: Thursday, September 3, Due Date: Thursday, September 17

**Instructions:** (1) All problem numbers below refer to the 4<sup>th</sup> edition of the textbook by Coluris, Dollimore and Kindberg. (2) Please, hand in hardcopy solutions that are typed (you may use your favorite word processor). We will not accept handwritten solutions. Figures and equations (if any) may be drawn by hand. (3) Please, start each problem on a fresh sheet and type your name at the top of each sheet. (4) Homework will be due at the beginning of class on the day of the deadline.

## Relevant Reading for this Homework: Sections 11.1-11.5

1. **(10 Points)** At 10:27:540 (hr, min, 1/100 sec.), server B requests time from the time-server  
A. At 10:27:610, server B receives a reply from timeserver A with the timestamp of 10:27:375.
  - a. **(5 Points)** Find out the drift of B's clock with respect to the time-server A's clock (assume there is no processing time at the time-server for time service).

**Answer 1a:** RTT: Reply – Request = 610-540 = 70 1/100sec.

Adjusted local time: Server + RTT/2 = 375 + 35 = 410 1/100sec.

Drift: Adjusted local time – local time = 410 – 610 = -200 1/100sec. = -2 sec

- b. **(2 Points)** Is B's clock going too fast or too slow? If the answer is yes, by how much is the clock going too fast or too slow?

**Answer 1b:** B's block is running 2 seconds **too fast**.

- c. **(3 Points)** How should B adjust its clock?

**Answer 1c:** To adjust B's clock we cannot just set the time back 2 seconds or we lose our monotonicity condition (The condition that the clock always moves only forward), which could result in events appearing as if they occurred in the future and a variety of issues in software relating to this problem. For B to adjust its clock safely, it should decrease the rate at which updates are given to its clock (at least in software) thus "slowing down" the time until it is caught up with the server's time.

2. **(5 Points)** In the symmetric mode of synchronization in NTP, suppose you are given that server A and server B are connected by a **symmetric channel**, i.e., the channel shows the same (but unknown) message delay both ways, i.e., the A to B delays is the same as B to A

\* late

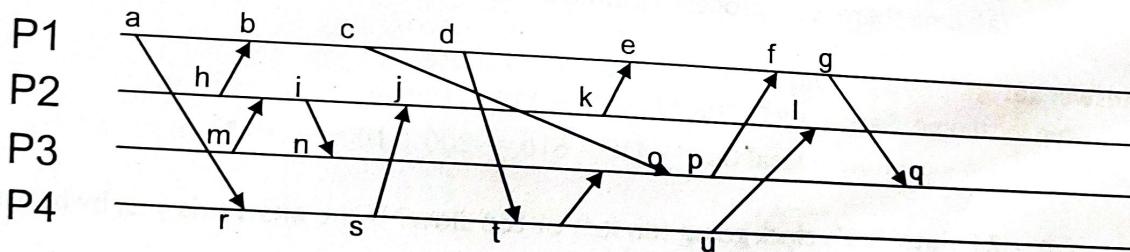
delay. Show using the equations for analyzing the NTP protocol, that under this situation, one can estimate the clock skew accurately (i.e., with an error of 0).

**Answer:** If the channel is symmetric, then  $t = t'$ , and

$$\Delta = \Delta + (t' - t)/2 = 0.$$

- That is, the actual offset (clock skew) is perfectly estimated (i.e., the error  $(t' - t)/2 = 0$ ).
3. (20 Points) Consider Figure 1 that shows four processes ( $P_1, P_2, P_3, P_4$ ) with events  $a, b, c, d, \dots$  all initialized to 0.

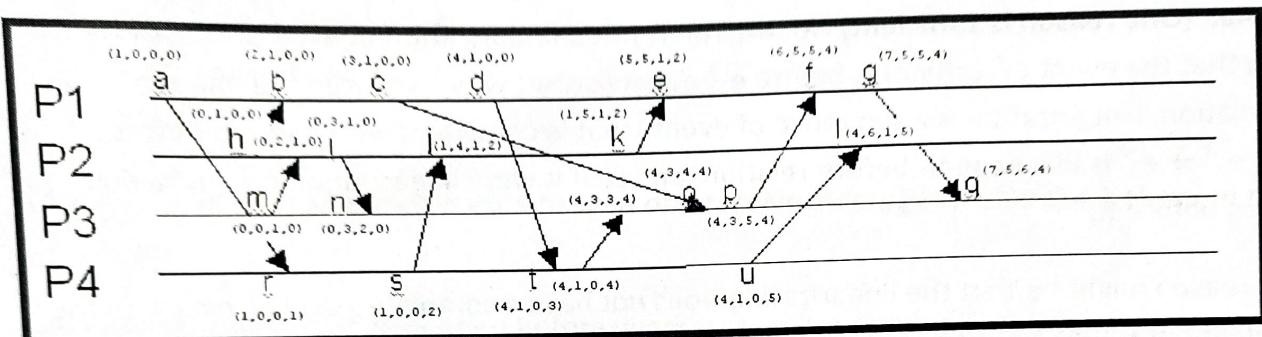
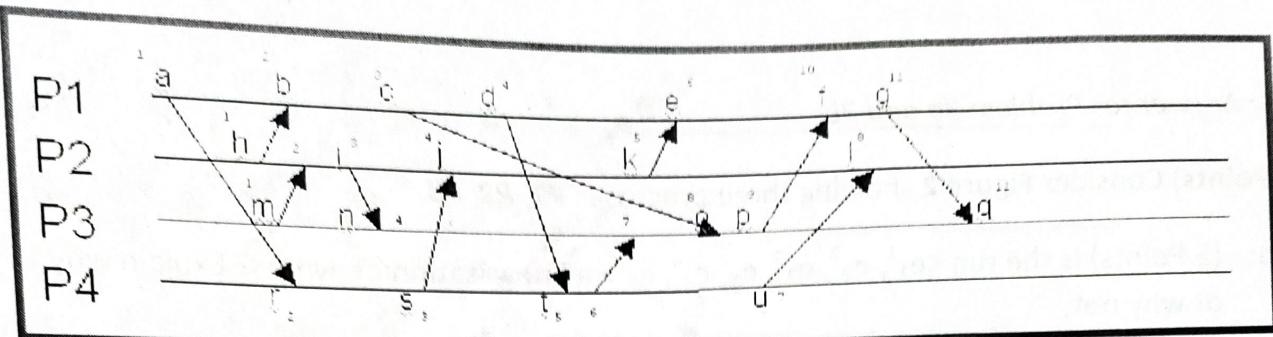
- a. (5 Points) List the **Lamport timestamps** for each event shown in Figure 1. Assume that each process maintains a logical clock as a single integer value as a Lamport clock. Provide timestamps for each labeled event.
- b. (10 Points) List the **Vector Clock timestamps** for each event shown in Figure 1. Provide timestamps for each labeled event.
- c. (5 Points) Is there the potential for a causal violation? Explain why.



**Figure 1:** Four Processes  $P_1, P_2, P_3, P_4$  run events  $a, b, c, d, \dots$  to send and receive messages

**Answer 3a&3b:** Answers for 3a and 3b are in **Figure 3**.

**Answer 3c:** Yes, there is a potential causality violation. Specifically from point C to point O, the message passed holds  $(3, 1, 0, 0)$  as the vector timestamp but the local vector timestamp on message arrival at point O is  $(4, 3, 3, 4)$ . Because the vector timestamp passed in the message is less than the local vector timestamp on arrival there is potential for a causal violation.



Event	Lamport Time Stamp	Vector Time Stamp
a	1	(1,0,0,0)
b	2	(2,1,0,0)
c	3	(3,1,0,0)
d	4	(4,1,0,0)
e	6	(5,5,1,2)
f	10	(6,5,5,4)
g	11	(7,5,5,4)
h	1	(0,1,0,0)
i	3	(0,3,1,0)
j	4	(1,4,1,2)
k	5	(1,5,1,2)
l	8	(4,6,1,5)
m	1	(0,0,1,0)
n	4	(0,3,2,0)
o	8	(4,3,4,4)
p	9	(4,3,5,4)
q	12	(7,5,6,4)
r	2	(1,0,0,1)
s	3	(1,0,0,2)
t	5	(4,1,0,3)
u	7	(4,1,0,5)

Figure 3: Answer for Problem 3a and 3b.

4. (35 Points) Consider Figure 2 showing three processes  $P1, P2, P3$ .

- a. (5 Points) Is the run  $\langle e_1^1, e_2^2, e_2^3, e_3^2, e_2^4, e_3^3 \rangle$  a linearization of events? Explain why or why not.

**Answer 4a:** (One reason is sufficient) No, the run is not a linearization of events. One of the reason is that the event  $e_2^4$  (effect) is before  $e_3^3$  event (cause) which violated the happen-before relation. Linearization is a sequence of events that is consistent with happen-before relation.  $e_3^3 \rightarrow e_2^4$  is in a happen-before relation, and so if it were linearization run,  $e_2^4$  would have been before  $e_3^3$ .

Another reason might be that the linearization does not have a complete set of all events in the global history of Figure 2.

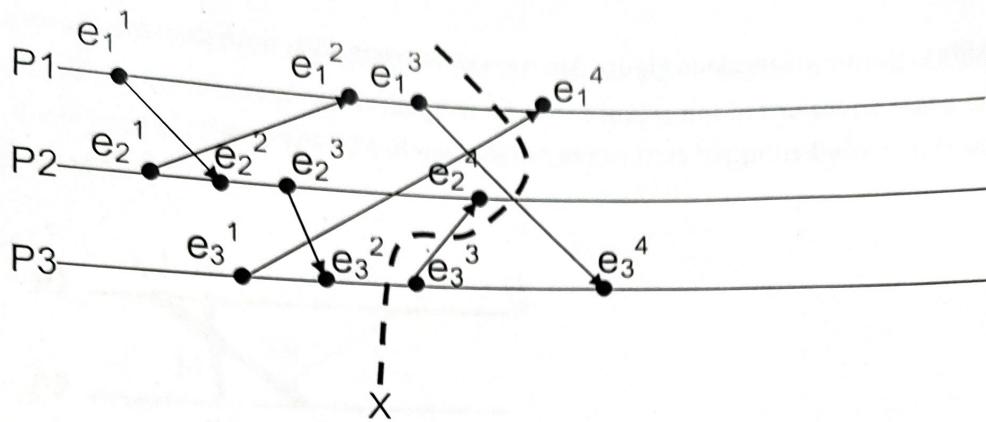
- b. (5 Points) Is the cut, shown by curve X, a consistent cut? Why?

**Answer 4b:** No. For a cut to be consistent, if  $e_i$  is in the cut and  $e_j$  happens before  $e_i$ , then  $e_j$  must be in the cut. In this example,  $e_2^4$  (receive event) is in the cut, but the corresponding send event ( $e_3^3$ ) that happens before it is not in the cut.

- c. (10 Points) Determine two other consistent cuts in Figure 2 and specify their frontiers.

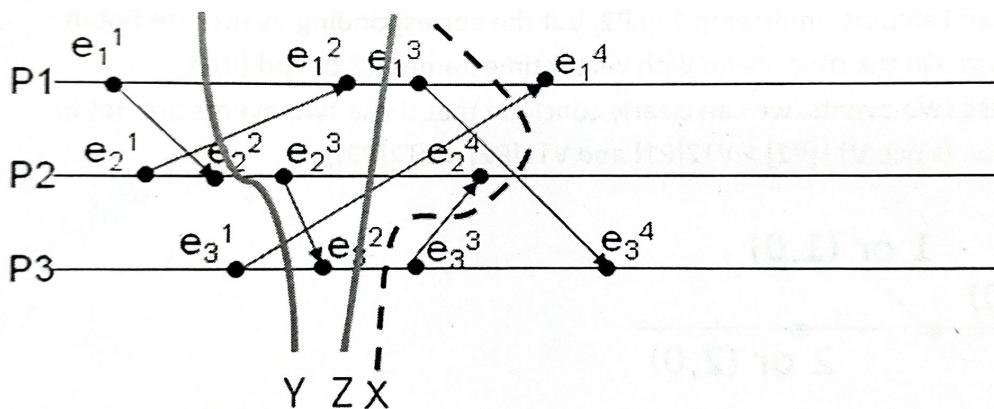
**Answer 4c:** See Figure 4 for two additional consistent cuts. Cuts Y and Z are consistent cuts. The frontier for Cut Y is  $\langle e_1^1, e_2^2, e_3^1 \rangle$ , The frontier for Cut Z is  $\langle e_1^2, e_2^3, e_3^2 \rangle$ .

- d. (15 Points) In the Figure 2, determine all the events that happen before event  $e_2^4$  (as per Lamport's Happened-Before relation). Also determine the events that are concurrent with event  $e_2^4$ .



**Figure 2:** Three Processes run events to send and receive messages with a Cut X through the processes.

**Answer 4d:** Events that happened before event  $e_2^4$ :  $e_1^1, e_2^1, e_2^2, e_2^3, e_3^1, e_3^2, e_3^3$ .  
 Events that are concurrent with event  $e_2^4$ :  $e_1^2, e_1^3, e_1^4, e_3^4$ .



**Figure 4:** Answer to Problem 4c – two additional consistent cuts Y and Z.

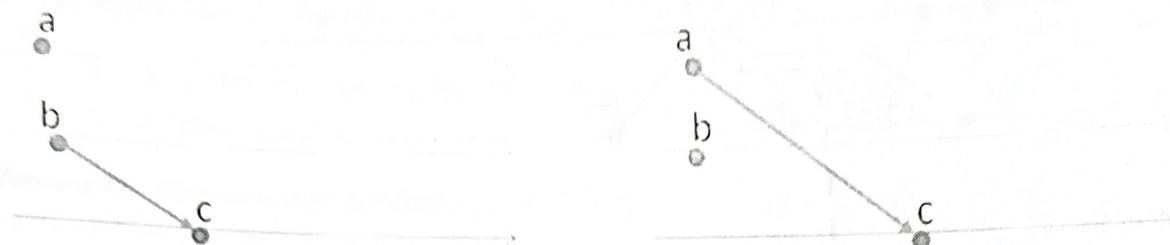
5. (10 Points)  $a, b$ , and  $c$  are events and no two events belong to the same process. Prove or disprove (give counter-example) the following:

- a. (5 Points)  $a$  is concurrent with  $b$  and  $b$  is before  $c$  implies that  $a$  is before  $c$ .

**Answer 5a:** No, consider counter example in **Figure 5a**:  $a$  and  $b$  are concurrent,  $b \rightarrow c$ ; however  $a$  is also concurrent with  $c$ .

- b. (5 Points)  $a$  is concurrent with  $b$  and  $b$  is concurrent with  $c$  implies that  $a$  is concurrent with  $c$ .

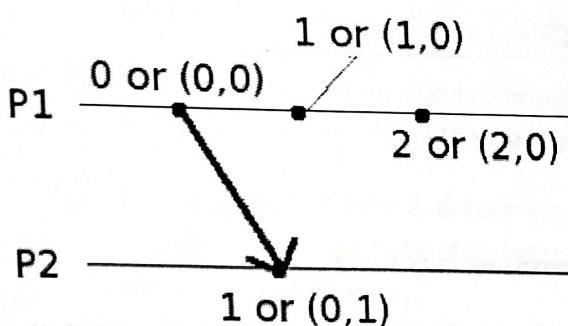
**Answer 5b:** No, consider counter example in Figure 5b:  $b$  is concurrent with both  $a$  and  $c$ , however,  $a \rightarrow c$ .



**Figure 5 : Counter Example for 5a (left), and Counter Example for 5b (right)**

6. **(10 Points)** Give an example where the Lamport's clock algorithm comes short (i.e., the Lamport's algorithm cannot clearly conclude that event  $e$  happens before  $e'$ , even those  $L(e) < L(e')$ , where  $L(e)$  is the Lamport's timestamp of the event  $e$ ), and the vector clock algorithm concludes clearly that event  $e$  happened before  $e'$  or NOT.

**Answer:** Consider the following example of 2 processes (Figure 6). The Lamport timestamp 2 in P1 is clearly larger than Lamport timestamp 1 in P2, but the corresponding events are not in happen-before relation. On the other hand with vector timestamps  $(2,0)$  and  $(0,1)$ , corresponding to those two events, we can clearly conclude that these two events are not in happen-before relation (since  $VT1[P1] > VT2[P1]$  and  $VT1[P2] < VT2[P2]$ ).



**Figure 6: Answer to Problem 6.**

7. **(10 Points)** If the FIFO channel assumption in the Chandy-Lamport algorithm is violated, then which step of the proof for the Chandy-Lamport algorithm given a consistent cut, breaks down?

**Answer:** The step of the proof that falls apart is: "If  $e_j$  occurred before  $p_j$  recorded its state, then  $e_i$  must have occurred before  $p_i$  recorded its state". Consider the Figure 7 with P1 sending generic message **M** and snapshot marker message **SS**. If the FIFO channel assumption is

violated, then it is possible that a message receive event (e3) is in the saved state of P2, even though the corresponding message send event (e2) is not in the saved state of P1. The resulting cut will contain the event e3 but exclude an event that happens-before it, thus is not consistent.

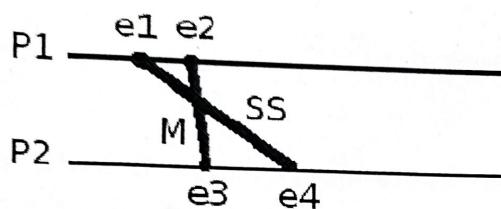


Figure 7: Answer to Problem 7.

# CSSE 550 Distributed Computing

## Mid-Term Exam #1

Name: \_\_\_\_\_

Note: All questions are based on textbook and Lamport's paper on "Time, Clocks, and the Ordering of Events in a Distributed System" to answer some of the following questions. **32 points in total. Limits your answers or solutions to at most 5 sentences.**

### Part I: True or False.

1. In Berkeley UNIX algorithm, one daemon asks all machines for their time and computes an average time and broadcasts this average time. True or False? [2 points]  
FALSE (broadcast adjustment)
2. The bully algorithm for coordinator election assumes that no process can fail during the election procedure until the new coordinator has been selected. True or False? [2 points]

False

**Part II: Short questions: keep your answers concise, complete and clean. Try to limit your answer to at most 5 sentences per question.**

3. What is Lamport's definition of *concurrent events* using the *happened before* relation? [4 points]

Given events a and b, a [not  $\rightarrow$ ] b, and b [not  $\rightarrow$ ] a

4. What is the advantage of a vector clock over Lamport's logical clock? [2 points]

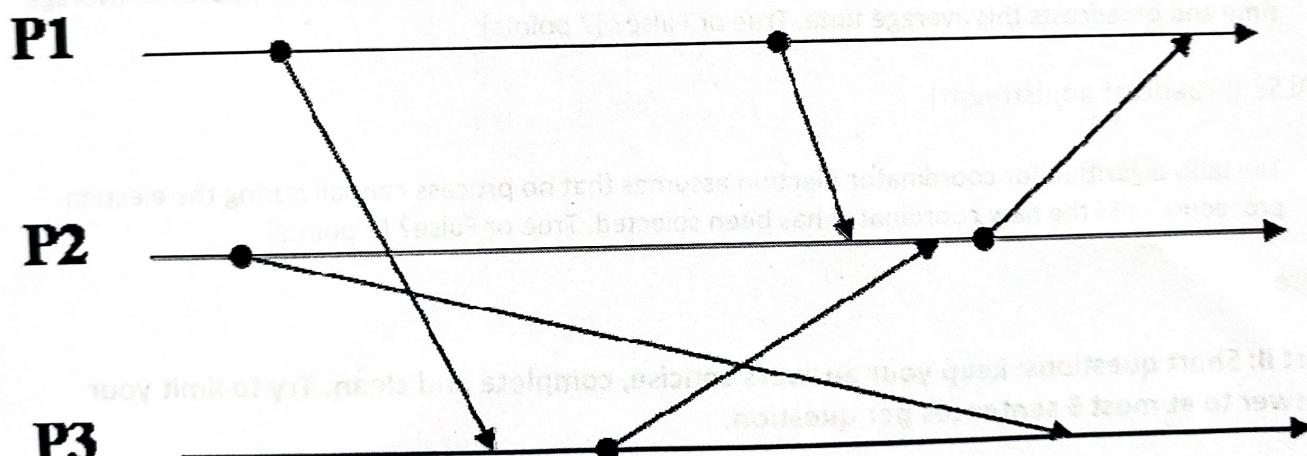
Catch causality

5. How do vector clocks extend the notion of Lamport's logical clocks? [4 points]

A vector clock maintains logical clocks of all involved processes. When a process  $P_i$  sends a message, it increases its logical clock in the vector by 1 and attach the updated logical clock to the message; Upon receiving a message from  $P_k$  with a vector clock  $VC(k)$ , for each component in  $P_i$ 's vector clock  $VC(i)$ , take the greater value in  $VC(k)$  and  $VC(i)$ . Optionally,  $P_i$  can increase its logical clock by 1.

Given  $VC(a) < VC(b)$  where a and b are events, we conclude  $a \rightarrow b$

6. Consider the following event diagram for processes P1, P2 and P3 executing in a distributed system. Compute the vector clock that is carried on each message. [10 points]



Solution 1:

$$\begin{aligned} P1 \rightarrow P3: & (1, 0, 0) \\ P2 \rightarrow P3: & (0, 1, 0) \\ P3 \rightarrow P2: & (1, 0, 1) \\ P1 \rightarrow P2: & (2, 0, 0) \\ P2 \rightarrow P1: & (2, 2, 1) \end{aligned}$$

Solution 2:

$$\begin{aligned} P1 \rightarrow P3: & (1, 0, 0) \\ P2 \rightarrow P3: & (0, 1, 0) \\ P3 \rightarrow P2: & (1, 0, 2) \\ P1 \rightarrow P2: & (2, 0, 0) \\ P2 \rightarrow P1: & (2, 4, 2) \end{aligned}$$

7. Suppose that two processes detect the demise of the coordinator simultaneously and both decide to hold an election using the bully algorithm. What happens? [4 points]

Each of the higher-numbered processes will get two *ELECTION* messages, but will ignore the second one. The election will proceed as usual.

8. You are synchronizing your clock from a time server using NTP and observe the following times: [4 points]
- a. timestamp at client when the message leaves the client: 6:22:15.100
  - b. timestamp at which the server receives the message: 7:05:10.700
  - c. timestamp at which the server sends the reply: 7:05:10.710
  - d. timestamp at client when the message is received at client: 6:22:15.250
- To what value do you set the client's clock?

$$T_2 - T_1 = 0:42:55.600$$

$$T_3 - T_4 = 0:42:55.460$$

$$\text{Theta} = [(T_2 - T_1) + (T_3 - T_4)] / 2 = 0:42:55.530$$

$\text{Theta} > 0$ , so we need to move the client's clock forward by 42 min 55 second and 530 msec  
the client's clock has to be set at time **6:22:15.250 + 00:42:55.530 = 07:05:10.780**